

Chap 13.

Transaction Processing

Concepts and Theory


DATABASE(CSI3105-01)

Donghyun Kang

Modified from Wonsuk Lee's Lecture notes
Images from Fundamentals of Database Systems 7th

Single-user vs Multi-user Systems

- Single-user DBMS : at most one user at a time can use DB system
- Multi-user DBMS : many users can use DB system concurrently
- single CPU : multi-programming OS
 - > interleaved execution (**interleaved concurrency**)
- multiple CPU : parallel processing(simultaneous concurrency)




Transactions

- Transaction :
 - execution of a program that forms a logical unit of database processing.
 - one or more atomic read/write operations
 - **read_item(X).**
Reads a database item named X into a program variable.
 - **write_item(X).**
Writes the value of program variable into the database item named X
- **read-only transaction**
- **read-write transaction**

Transactions

- Two transactions T1 and T2 are requested at the same time
 - correct result of interleaved execution
 - => either T1 -> T2 or T2 -> T1



Transactions

- **Sample transaction**

```
1      EXEC SQL WHENEVER SQLERROR GOTO UNDO:
2      EXEC SQL SET TRANSACTION
          READ WRITE
          DIAGNOSTIC SIZE 5
          ISOLATION LEVEL SERIALIZABLE;
3      EXEC SQL INSERT INTO EMPLOYEE (Fname,Lname,Ssn,Dno,Salary)
          VALUES('robert','smith','991004321',2,35000);
4      EXEC SQL UPDATE EMPLOYEE
          SET salary = salary *1.1 WHERE dno=2;
5      EXEC SQL COMMIT;
6      GOTO THE_END:
7 UNDO:    EXEC SQL ROLLBACK;
8 THE_END: ....;
```

Concurrency Control

- Why Concurrency Control?

[Example] Airline reservation system

- T1 : cancels N reservations in X flight & reserves N seats in Y flight
- T2 : reserves M seats in X flight

initial value of
X= 80,
N=5,
M=4

(a)

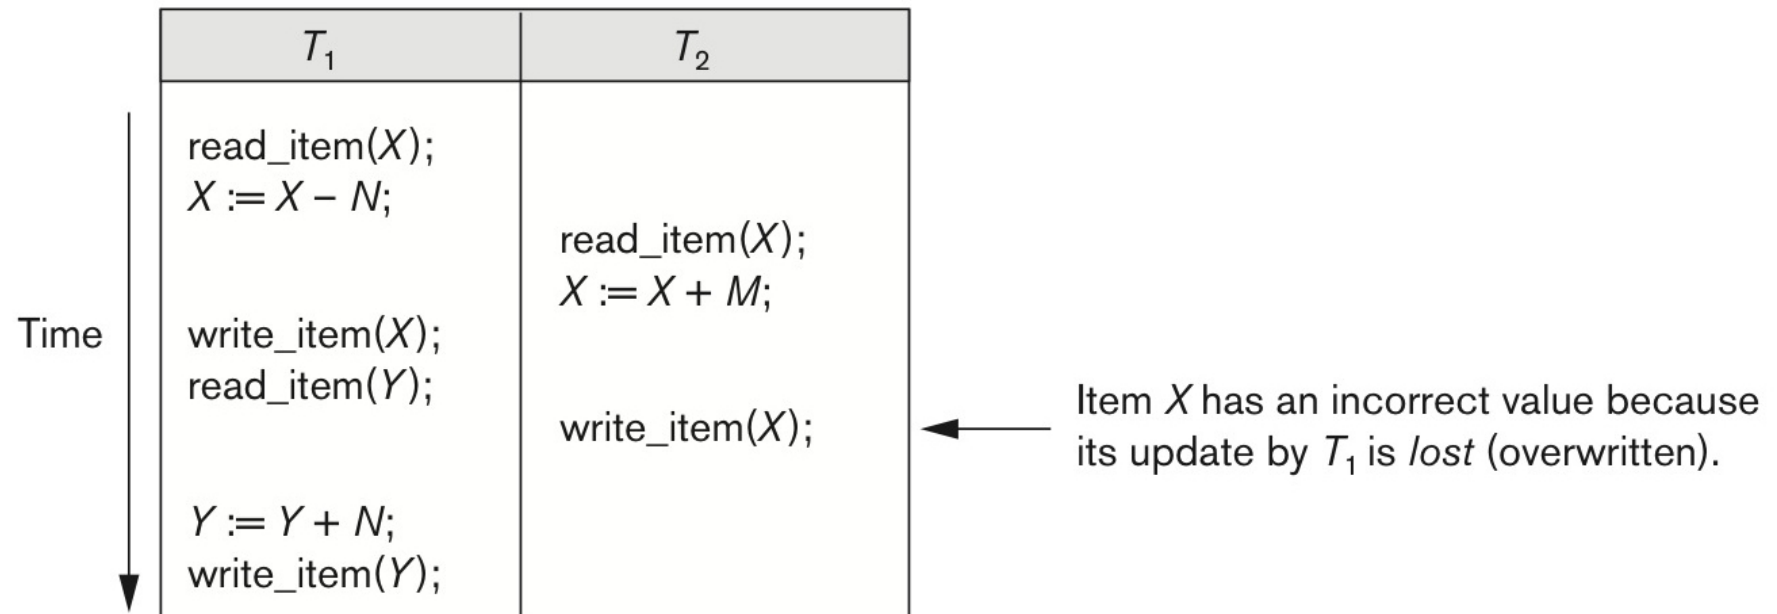
T_1
<pre>read_item(X); X := X - N; write_item(X); read_item(Y); Y := Y + N; write_item(Y);</pre>

(b)

T_2
<pre>read_item(X); X := X + M; write_item(X);</pre>

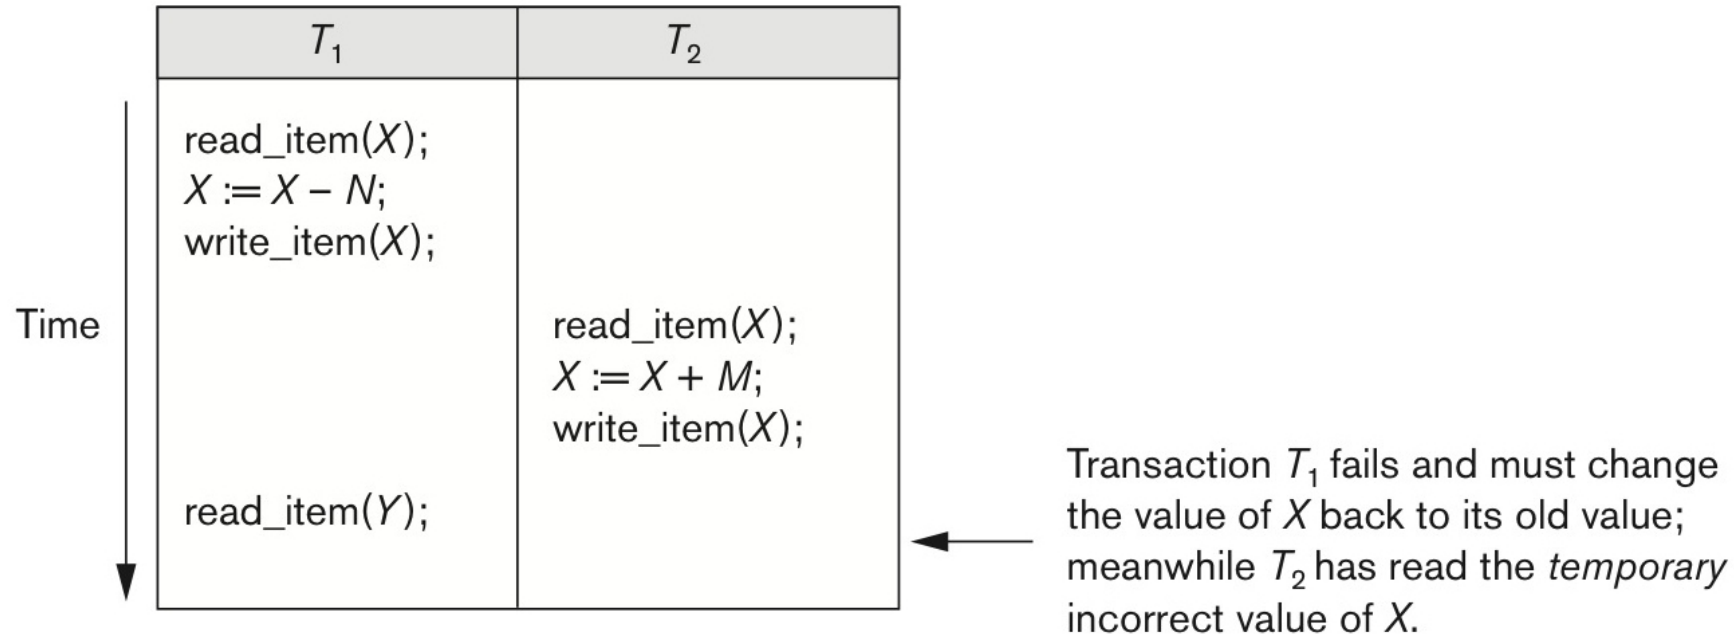
lost update problem

- two transactions that access the same database items have their operations interleaved in a way that makes the value of some database items incorrect



Temporary update problem

- One transaction updates a DB item and fail for some reason
 - The updated item is accessed by another transaction before it is changed back to its original value



The Incorrect Summary Problem

- one transaction is calculating an aggregate summary function on a number of database items
 - while other transactions are updating some of these items

T_1	T_3
<pre>read_item(X); X := X - N; write_item(X); read_item(Y); Y := Y + N; write_item(Y);</pre>	<pre>sum := 0; read_item(A); sum := sum + A; ⋮ read_item(X); sum := sum + X; read_item(Y); sum := sum + Y;</pre>

← T_3 reads X after N is subtracted and reads Y before N is added; a wrong summary is the result (off by N).

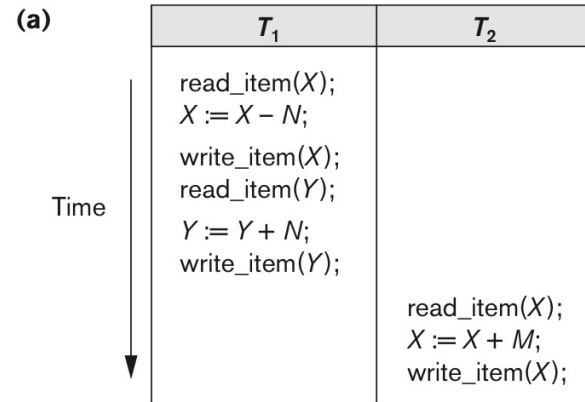
The Unrepeatable Read Problem

- a transaction T1 reads the same item twice
 - the item is changed by another transaction T2 between the two reads.
 - Hence, T1 receives different values for its two reads of the same item.

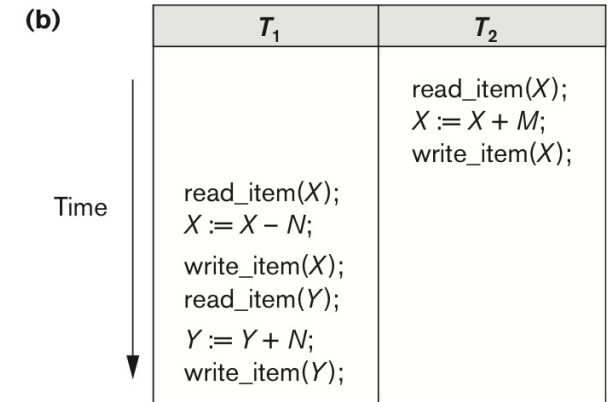
T1	T2
Read_item(X)	X:=X*1.1 Write_item(X)
Read_item(X)	

Schedule & serializable

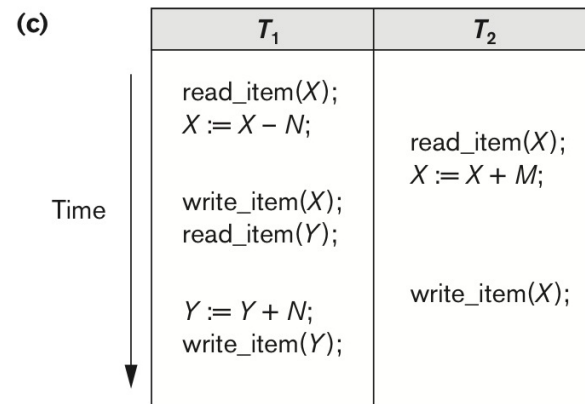
- A **schedule** S of n transactions is an ordering of the operations of the transaction
- Schedule S is **serial** if, for every transaction T participating in the schedule, all the operations of T are executed consecutively in the schedule
- Schedule S is **Serializable** if it is *equivalent* to some serial schedule of the same n transactions.



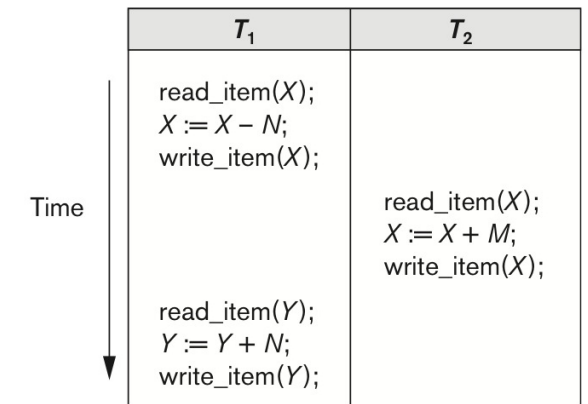
Schedule A



Schedule B



Schedule C



Schedule D

Schedule & serializable

- Two schedules are Equivalent
 - Result equivalent
 - the same final state of the database.
 - View equivalent
 - as long as each read operation of a transaction reads the result of the same write operation in both schedules, the write operations of each transaction must produce the same results
- **Conflict equivalent**
 - if the relative order of any two *conflicting operations* is the same in both schedules.
 - two operations in a schedule are *conflict* if they belong to different transactions, access the same database item, and either both are write_item operations or one is a write_item and the other a read_item

TWO Phase Locking (2PL)

- Two Phase Locking
 - multiple-mode locks (read & write locks) + protocols to guarantee serializability
 - Transaction T follows 2PL if all locking operations precede the first unlock operation in T
- (i) expanding phase : acquire new locks
(no unlock op)
- (ii) shrinking phase : unlock all locks (no lock op)

TWO Phase Locking (2PL)

(a)

T_1	T_2
<pre>read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X);</pre>	<pre>read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y);</pre>

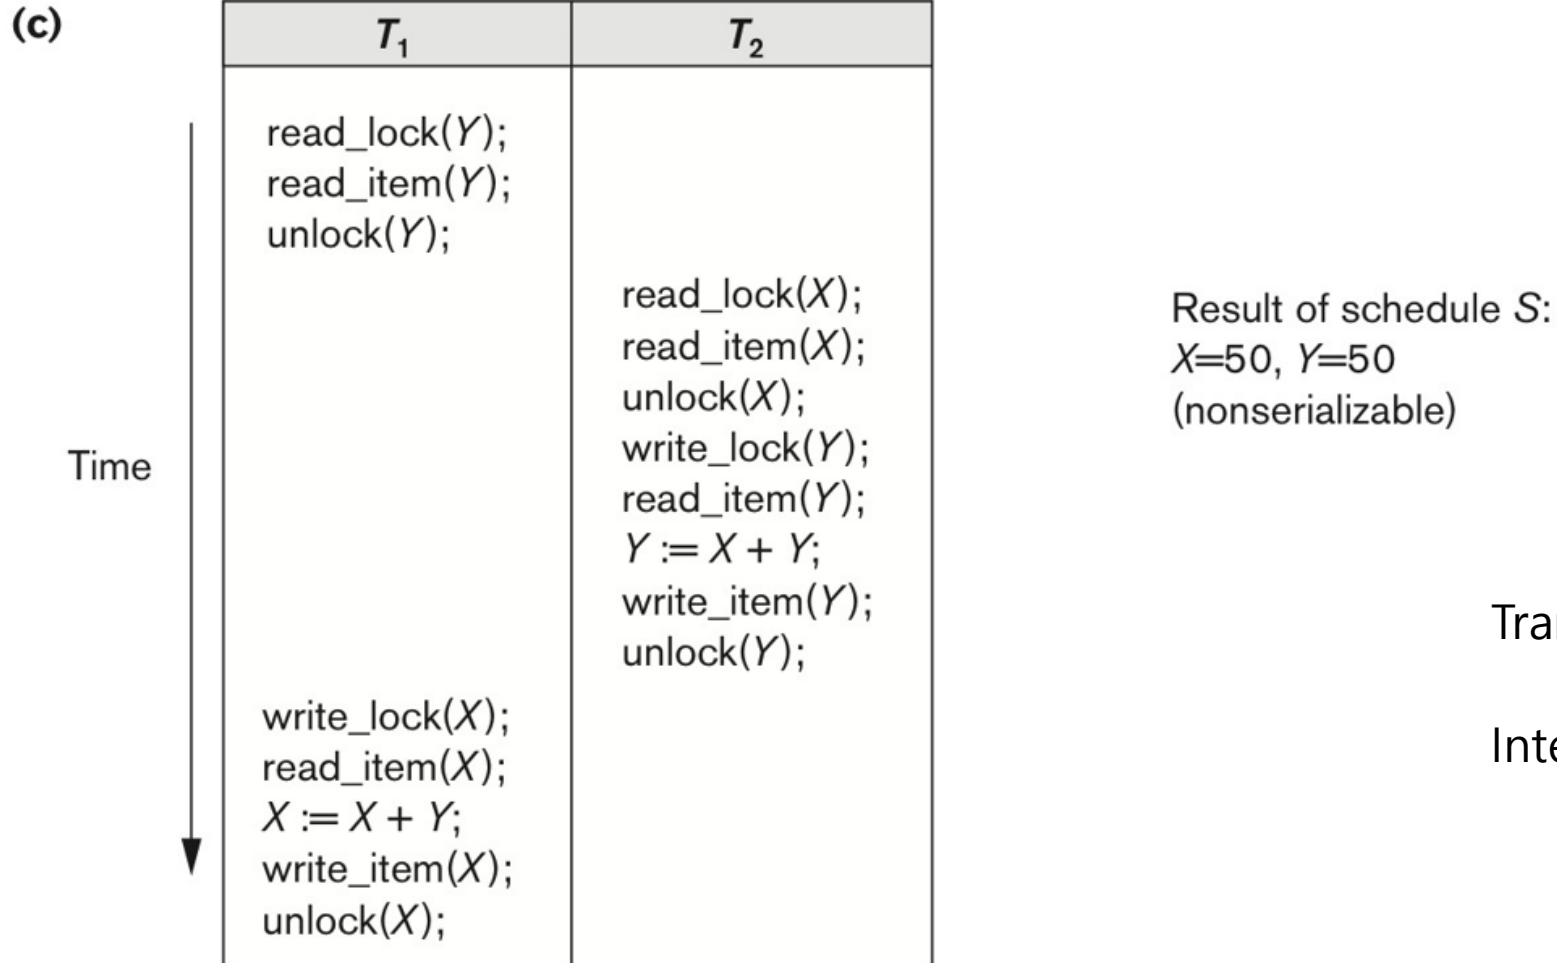
(b)

Initial values: $X=20, Y=30$

Result serial schedule T_1
followed by T_2 : $X=50, Y=80$

Result of serial schedule T_2
followed by T_1 : $X=70, Y=50$

TWO Phase Locking (2PL)



Transactions that do not obey 2PL.

Interleaved execution of T1 and T2.

TWO Phase Locking (2PL)

T_1'
read_lock(Y); read_item(Y); write_lock(X); unlock(Y) read_item(X); $X := X + Y$; write_item(X); unlock(X);

T_2'
read_lock(X); read_item(X); write_lock(Y); unlock(X) read_item(Y); $Y := X + Y$; write_item(Y); unlock(Y);

수정사항

write lock(X)

$X = X + X$

unlock(X)

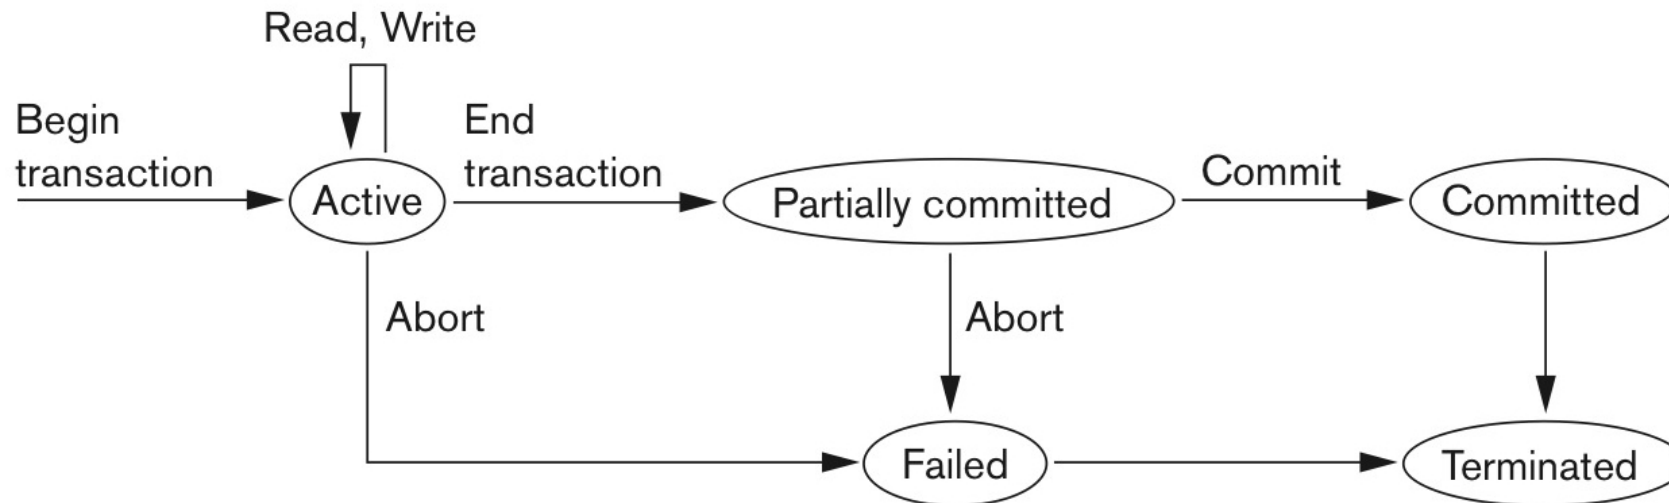
Transactions that follow 2PL

Recovery

- Why Recovery?
 - transaction should be atomic
 - all(**committed**) or nothing(**aborted**)
- Types of Failures
 - computer failure(**system crash**)
 - transaction or system error
 - Local errors or exception conditions detected by the transaction
 - Concurrency control enforcement.
 - Disk failure.
 - Physical problems and catastrophes

Recovery

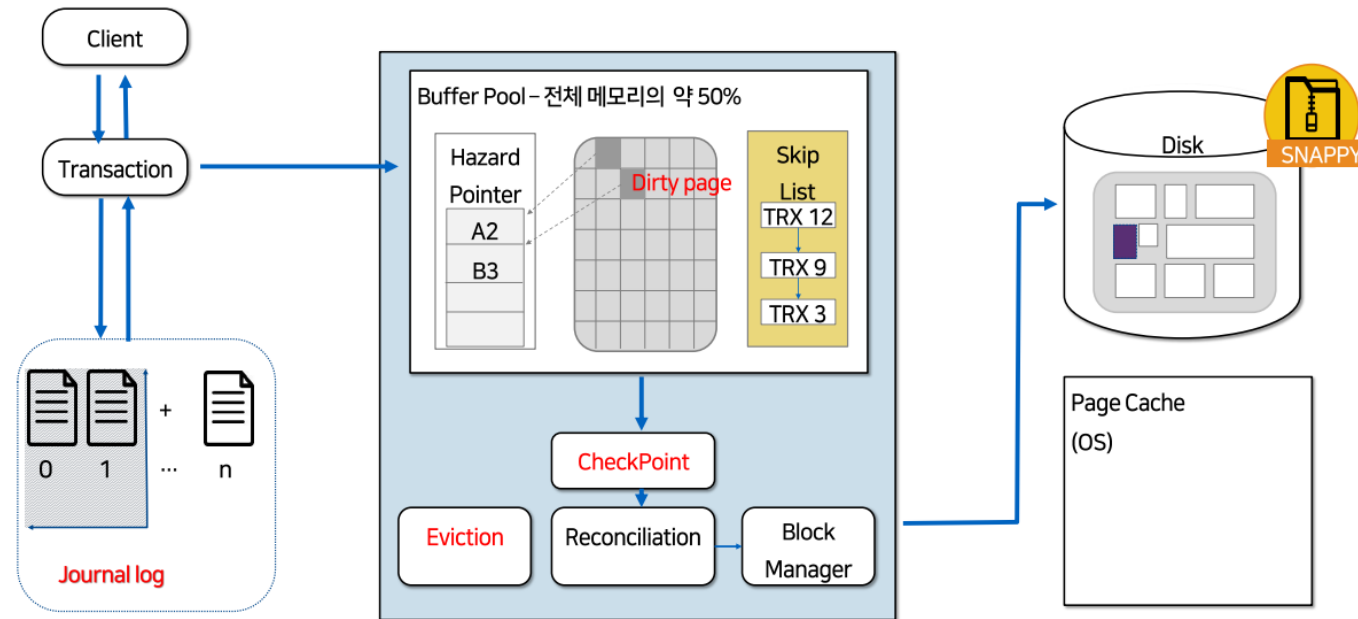
- Committed transaction
 - a transaction which is executed successfully.
 - recovery manager only cares for committed transactions
- commit point
 - commit record [commit, transaction T] into the log.



Recovery

DEVVIEW
2018

3.6 & 7 eviction & Checkpoint



- **checkpoint** : [checkpoint, list of active transactions] log
 - the system writes out to the database on disk all DBMS buffers that have been modified.
 - all operations in a transaction have been executed successfully and recorded in the log file [commit, T]

Recovery

- Recovery:
 - system fail => return to consistent DB state
- **if** catastrophic failure,
 - tape dump & log file
 - restructuring old DB state to current state
by redoing committed transaction
- **else if** inconsistent DB state
 - log file
 - undo & redo (-> consistent DB state)

Recovery

(i) deferred update

: after commit point, all updates by a transaction are recorded persistently

=> change the content of DB in disk after commit point


=> No-Undo/Redo algorithm

(ii) immediate update

: update DB in disk before commit point

=> Undo/Redo algorithm

Recovery



Properties of Transactions (ACID)

- Desirable Properties of Transactions
 - (i) **A**tomicity : transaction = atomic unit of processing (recovery manager)
 - (ii) **C**onsistency perservation : (programmer)
consistent DB state -> another consistent DB state
 - (iii) **I**solation : all updates in a transaction should not be visible to other transactions until it is committed (concurrency control)
 - (iv) **D**urability (permanency) : (recovery manager)
 - Once a transaction is committed (changes DB), the changes should not be lost because of subsequent failure